

[🏠](#) > [User Guide](#) > [Chart visualization](#)

Chart visualization

Note

The examples below assume that you're using [Jupyter](#).

This section demonstrates visualization through charting. For information on visualization of tabular data please see the section on [Table Visualization](#).

We use the standard convention for referencing the matplotlib API:

```
In [1]: import matplotlib.pyplot as plt
```

```
In [2]: plt.close("all")
```

We provide the basics in pandas to easily create decent looking plots. See [the ecosystem page](#) for visualization libraries that go beyond the basics documented here.

[Skip to main content](#)

Note

All calls to `np.random` are seeded with 123456.

Basic plotting: `plot`

We will demonstrate the basics, see the [cookbook](#) for some advanced strategies.

The `plot` method on Series and DataFrame is just a simple wrapper around `plt.plot()`:


```
In [3]: np.random.seed(123456)

In [4]: ts = pd.Series(np.random.randn(1000), index=pd.date_range("1/1/2000", periods=1000))

In [5]: ts = ts.cumsum()

In [6]: ts.plot();
```

[Skip to main content](#)




If the index consists of dates, it calls `gcf().autofmt_xdate()` to try to format the x-axis nicely as per above.

On DataFrame, `plot()` is a convenience to plot all of the columns with labels:

```
In [7]: df = pd.DataFrame(np.random.randn(1000, 4), index=ts.index, columns=list("ABCD"))  
In [8]: df = df.cumsum()  
In [9]: plt.figure();
```

[Skip to main content](#)

```
In [10]: df.plot();
```




You can plot one column versus another using the `x` and `y` keywords in `plot()`:

```
In [11]: df3 = pd.DataFrame(np.random.randn(1000, 2), columns=["B", "C"]).cumsum()
```

```
In [12]: df3["A"] = pd.Series(list(range(len(df))))
```

```
In [13]: df3.plot(x="A", y="B")
```

[Skip to main content](#)



Note

For more formatting and styling options, see [formatting](#) below.

Other plots

[Skip to main content](#)


Plotting methods allow for a handful of plot styles other than the default line plot. These methods can be provided as the `kind` keyword argument to `plot()`, and include:

- `'bar'` or `'barh'` for bar plots
- `'hist'` for histogram
- `'box'` for boxplot
- `'kde'` or `'density'` for density plots
- `'area'` for area plots
- `'scatter'` for scatter plots
- `'hexbin'` for hexagonal bin plots
- `'pie'` for pie plots

For example, a bar plot can be created the following way:

```
In [14]: plt.figure();  
  
In [15]: df.iloc[5].plot(kind="bar");
```

[Skip to main content](#)



You can also create these other plots using the methods `DataFrame.plot.<kind>` instead of providing the `kind` keyword argument. This makes it easier to discover plot methods and the specific arguments they use:

```
In [16]: df = pd.DataFrame()
```

```
In [17]: df.plot.<TAB> # noqa: E225, E999
```

<code>df.plot.area</code>	<code>df.plot.barh</code>	<code>df.plot.density</code>	<code>df.plot.hist</code>	<code>df.plot.line</code>	<code>df.p</code>
<code>df.plot.bar</code>	<code>df.plot.box</code>	<code>df.plot.hexbin</code>	<code>df.plot.kde</code>	<code>df.plot.pie</code>	

[Skip to main content](#)

In addition to these `kind` s, there are the `DataFrame.hist()`, and `DataFrame.boxplot()` methods, which use a separate interface.

Finally, there are several plotting functions in `pandas.plotting` that take a `Series` or `DataFrame` as an argument. These include:

- Scatter Matrix
- Andrews Curves
- Parallel Coordinates
- Lag Plot
- Autocorrelation Plot
- Bootstrap Plot
- RadViz


Plots may also be adorned with `errorbars` or `tables`.

Bar plots

For labeled, non-time series data, you may wish to produce a bar plot:

```
In [18]: plt.figure();  
  
In [19]: df.iloc[5].plot.bar();  
  
In [20]: plt.axhline(0, color="k");
```

[Skip to main content](#)




Calling a DataFrame's `plot.bar()` method produces a multiple bar plot:

```
In [21]: df2 = pd.DataFrame(np.random.rand(10, 4), columns=["a", "b", "c", "d"])
```

```
In [22]: df2.plot.bar();
```


[Skip to main content](#)



To produce a stacked bar plot, pass `stacked=True`:

```
In [23]: df2.plot.bar(stacked=True);
```


[Skip to main content](#)



To get horizontal bar plots, use the `barh` method:

```
In [24]: df2.plot.barh(stacked=True);
```

[Skip to main content](#)



Histograms

Histograms can be drawn by using the `DataFrame.plot.hist()` and `Series.plot.hist()` methods


```
In [25]: df4 = pd.DataFrame(  
    ....:     {  
    ....:         "a": np.random.randn(1000) + 1,  
    ....:         "b": np.random.randn(1000)
```

[Skip to main content](#)

```
....:     columns=["a", "b", "c"],  
....: )  
....:
```

```
In [26]: plt.figure();
```

```
In [27]: df4.plot.hist(alpha=0.5);
```




A histogram can be stacked using `stacked=True`. Bin size can be changed using the `bins` keyword.

[Skip to main content](#)

```
In [28]: plt.figure();
```

```
In [29]: df4.plot.hist(stacked=True, bins=20);
```




You can pass other keywords supported by matplotlib `hist`. For example, horizontal and cumulative histograms can be drawn by `orientation='horizontal'` and `cumulative=True`.

```
In [30]: plt.figure();
```

[Skip to main content](#)

```
In [31]: df4["a"].plot.hist(orientation="horizontal", cumulative=True);
```




See the `hist` method and the [matplotlib hist documentation](#) for more.

The existing interface `DataFrame.hist` to plot histogram still can be used.

```
In [32]: plt.figure();
```

```
In [33]: df["A"].diff().hist();
```

[Skip to main content](#)




`DataFrame.hist()` plots the histograms of the columns on multiple subplots:

```
In [34]: plt.figure();
```

```
In [35]: df.diff().hist(color="k", alpha=0.5, bins=50);
```

[Skip to main content](#)




The `by` keyword can be specified to plot grouped histograms:

```
In [36]: data = pd.Series(np.random.randn(1000))
```

```
In [37]: data.hist(by=np.random.randint(0, 4, 1000), figsize=(6, 4));
```

[Skip to main content](#)



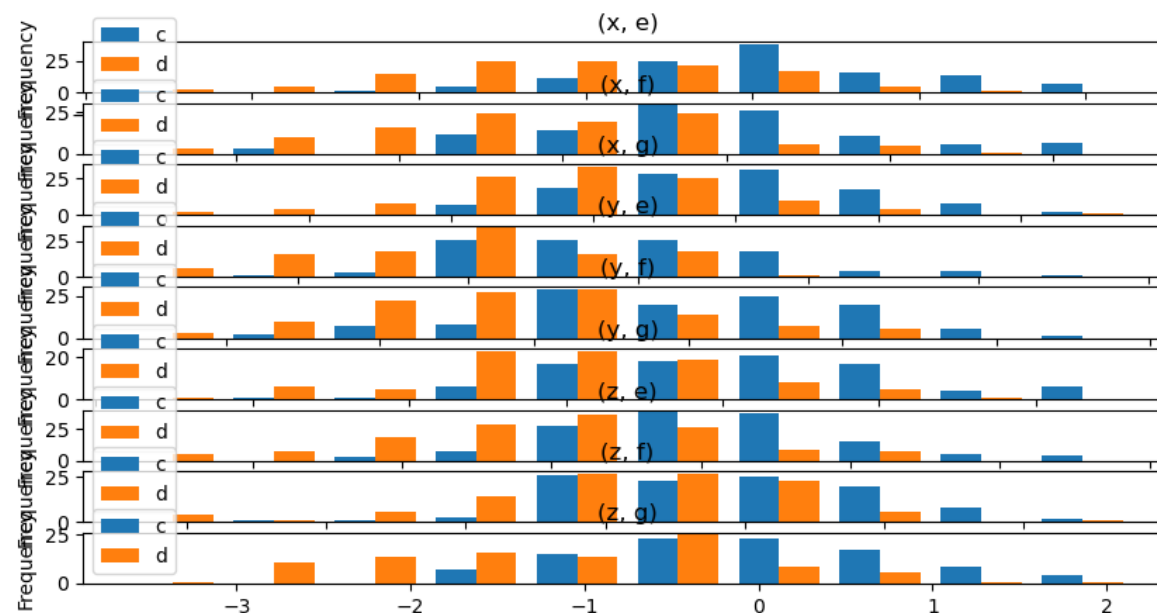
In addition, the `by` keyword can also be specified in `DataFrame.plot.hist()`.

! Changed in version 1.4.0.

```
In [38]: data = pd.DataFrame(
...:     {
...:         "a": np.random.choice(["x", "y", "z"], 1000),
...:         "b": np.random.choice(["e", "f", "g"], 1000),
...:         "c": np.random.randn(1000),
...:         "d": np.random.randn(1000) - 1,
...:     },
...: )
...:
```

[Skip to main content](#)

```
In [39]: data.plot.hist(by=["a", "b"], figsize=(10, 5));
```



Box plots


Boxplot can be drawn calling `Series.plot.box()` and `DataFrame.plot.box()`, or `DataFrame.boxplot()` to visualize the distribution of values within each column.

For instance, here is a boxplot representing five trials of 10 observations of a uniform random variable on $[0,1)$.

```
In [40]: df = pd.DataFrame(np.random.rand(10, 5), columns=["A", "B", "C", "D", "E"])
```

[Skip to main content](#)

```
In [41]: df.plot.box();
```



Boxplot can be colored by passing `color` keyword. You can pass a `dict` whose keys are `boxes`, `whiskers`, `medians` and `caps`. If some keys are missing in the `dict`, default colors are used for the corresponding artists. Also, boxplot has `sym` keyword to specify fliers style.


When you pass other type of arguments via `color` keyword, it will be directly passed to matplotlib for all the `boxes`, `whiskers`, `medians` and `caps` colorization.

[Skip to main content](#)

The colors are applied to every boxes to be drawn. If you want more complicated colorization, you can get each drawn artists by passing `return_type`.

```
In [42]: color = {  
.....:     "boxes": "DarkGreen",  
.....:     "whiskers": "DarkOrange",  
.....:     "medians": "DarkBlue",  
.....:     "caps": "Gray",  
.....: }  
.....:  
  
In [43]: df.plot.box(color=color, sym="r+");
```


[Skip to main content](#)



Also, you can pass other keywords supported by matplotlib `boxplot`. For example, horizontal and custom-positioned boxplot can be drawn by `vert=False` and `positions` keywords.

```
In [44]: df.plot.box(vert=False, positions=[1, 4, 5, 6, 8]);
```

[Skip to main content](#)



See the `boxplot` method and the [matplotlib boxplot documentation](#) for more.


The existing interface `DataFrame.boxplot` to plot boxplot still can be used.

```
In [45]: df = pd.DataFrame(np.random.rand(10, 5))
```

```
In [46]: plt.figure();
```

```
In [47]: bp = df.boxplot()
```


[Skip to main content](#)



You can create a stratified boxplot using the `by` keyword argument to create groupings. For instance,

```
In [48]: df = pd.DataFrame(np.random.rand(10, 2), columns=["Col1", "Col2"])
In [49]: df["X"] = pd.Series(["A", "A", "A", "A", "A", "B", "B", "B", "B", "B"])
In [50]: plt.figure();
In [51]: bp = df.boxplot(by="X")
```


[Skip to main content](#)



You can also pass a subset of columns to plot, as well as group by multiple columns:

```
In [52]: df = pd.DataFrame(np.random.rand(10, 3), columns=["Col1", "Col2", "Col3"])
In [53]: df["X"] = pd.Series(["A", "A", "A", "A", "A", "B", "B", "B", "B", "B"])
In [54]: df["Y"] = pd.Series(["A", "B", "A", "B", "A", "B", "A", "B", "A", "B"])
In [55]: plt.figure();
In [56]: bp = df.boxplot(column=["Col1", "Col2"], by=["X", "Y"])
```

[Skip to main content](#)



You could also create groupings with `DataFrame.plot.box()`, for instance:

! Changed in version 1.4.0.


```
In [57]: df = pd.DataFrame(np.random.rand(10, 3), columns=["Col1", "Col2", "Col3"])
```

```
In [58]: df["X"] = pd.Series(["A", "A", "A", "A", "A", "B", "B", "B", "B", "B"])
```

```
In [59]: df.plot.figure()
```

[Skip to main content](#)

```
In [60]: bp = df.plot.box(column=["Col1", "Col2"], by="X")
```



In `boxplot`, the return type can be controlled by the `return_type` keyword. The valid choices are `{"axes", "dict", "both", None}`. Faceting, created by `DataFrame.boxplot` with the `by` keyword, will affect the output type as well:

`return_type`

Faceted

Output type

[Skip to main content](#)

<code>return_type</code>	Faceted	Output type
<code>None</code>	Yes	2-D ndarray of axes
<code>'axes'</code>	No	axes
<code>'axes'</code>	Yes	Series of axes
<code>'dict'</code>	No	dict of artists
<code>'dict'</code>	Yes	Series of dicts of artists
<code>'both'</code>	No	namedtuple
<code>'both'</code>	Yes	Series of namedtuples

`Groupby.boxplot` always returns a `Series` of `return_type`.

```
In [61]: np.random.seed(1234)


In [62]: df_box = pd.DataFrame(np.random.randn(50, 2))

In [63]: df_box["g"] = np.random.choice(["A", "B"], size=50)

In [64]: df_box.loc[df_box["g"] == "B", 1] += 3

In [65]: bp = df_box.boxplot(by="g")
```

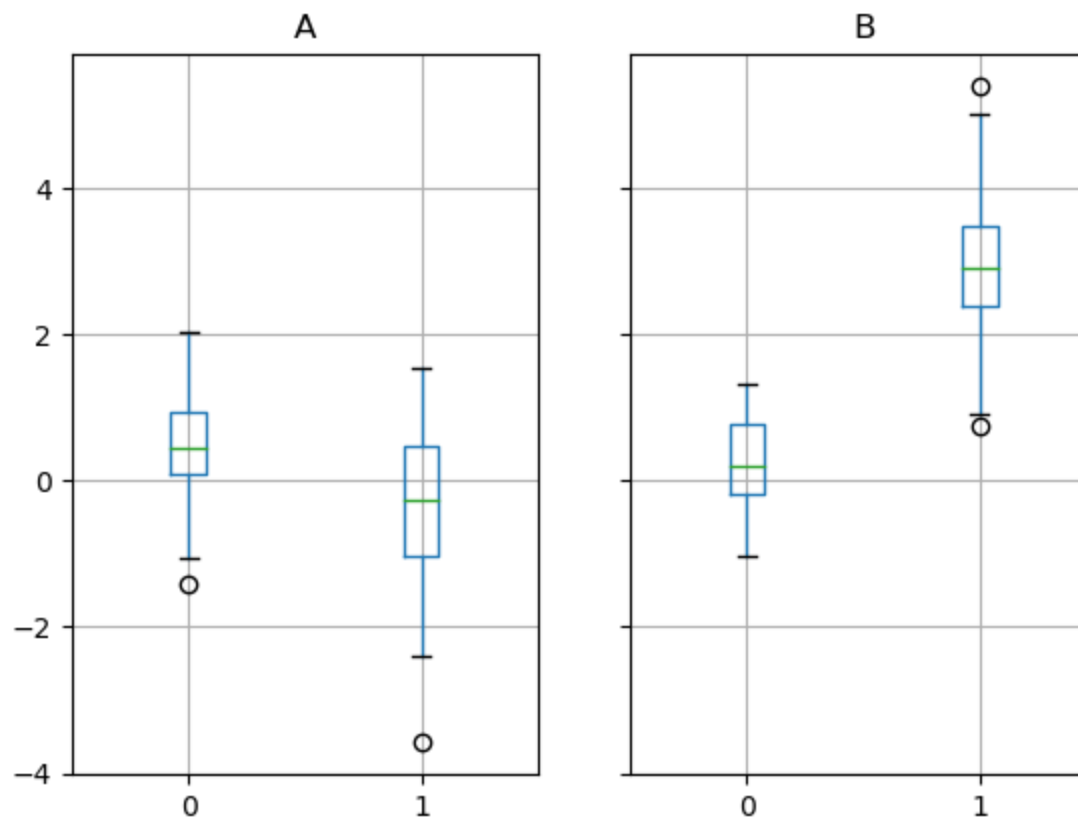
[Skip to main content](#)



The subplots above are split by the numeric columns first, then the value of the `g` column. Below the subplots are first split by the value of `g`, then by the numeric columns.

```
In [66]: bp = df_box.groupby("g").boxplot()
```

[Skip to main content](#)



Area plot


You can create area plots with `Series.plot.area()` and `DataFrame.plot.area()`. Area plots are stacked by default. To produce stacked area plot, each column must be either all positive or all negative values.

When input data contains `NaN`, it will be automatically filled by 0. If you want to drop or fill by

[Skip to main content](#)

```
In [67]: df = pd.DataFrame(np.random.rand(10, 4), columns=["a", "b", "c", "d"])
```


```
In [68]: df.plot.area();
```



To produce an unstacked plot, pass `stacked=False`. Alpha value is set to 0.5 unless otherwise specified:

```
In [69]: df.plot.area(stacked=False);
```

[Skip to main content](#)



Scatter plot


Scatter plot can be drawn by using the `DataFrame.plot.scatter()` method. Scatter plot requires numeric columns for the x and y axes. These can be specified by the `x` and `y` keywords.

```
In [70]: df = pd.DataFrame(np.random.rand(50, 4), columns=["a", "b", "c", "d"])
```

[Skip to main content](#)


```
....: )  
....:
```

```
In [72]: df.plot.scatter(x="a", y="b");
```




To plot multiple column groups in a single axes, repeat `plot` method specifying target `ax`. It is recommended to specify `color` and `label` keywords to distinguish each groups.

```
In [73]: ax = df.plot.scatter(x="a", y="b", color="DarkBlue", label="Group 1")
```

[Skip to main content](#)


```
In [74]: df.plot.scatter(x="c", y="d", color="DarkGreen", label="Group 2", ax=ax);
```



The keyword `c` may be given as the name of a column to provide colors for each point:

```
In [75]: df.plot.scatter(x="a", y="b", c="c", s=50);
```

[Skip to main content](#)




If a categorical column is passed to `c`, then a discrete colorbar will be produced:

! *New in version 1.3.0.*

```
In [76]: df.plot.scatter(x="a", y="b", c="species", cmap="viridis", s=50);
```


[Skip to main content](#)



You can pass other keywords supported by matplotlib `scatter`. The example below shows a bubble chart using a column of the `DataFrame` as the bubble size.

```
In [77]: df.plot.scatter(x="a", y="b", s=df["c"] * 200);
```

[Skip to main content](#)



See the `scatter` method and the [matplotlib scatter documentation](#) for more.

Hexagonal bin plot


You can create hexagonal bin plots with `DataFrame.plot.hexbin()`. Hexbin plots can be a useful alternative to scatter plots if your data are too dense to plot each point individually.

```
In [79]: df = pd.DataFrame(np.random.randn(1000, 2), columns=["a", "b"])
```

[Skip to main content](#)

```
In [79]: df["b"] = df["b"] + np.arange(1000)

In [80]: df.plot.hexbin(x="a", y="b", gridsize=25);
```



A useful keyword argument is `gridsize`; it controls the number of hexagons in the x-direction, and defaults to 100. A larger `gridsize` means more, smaller bins.

By default, a histogram of the counts around each `(x, y)` point is computed. You can specify alternative aggregations by passing values to the `C` and `reduce_C_function` arguments. `C`

[Skip to main content](#)

that reduces all the values in a bin to a single number (e.g. `mean`, `max`, `sum`, `std`). In this example the positions are given by columns `a` and `b`, while the value is given by column `z`. The bins are aggregated with NumPy's `max` function.


```
In [81]: df = pd.DataFrame(np.random.randn(1000, 2), columns=["a", "b"])
```

```
In [82]: df["b"] = df["b"] + np.arange(1000)
```

```
In [83]: df["z"] = np.random.uniform(0, 3, 1000)
```

```
In [84]: df.plot.hexbin(x="a", y="b", C="z", reduce_C_function=np.max, gridsize=25);
```

[Skip to main content](#)




See the `hexbin` method and the [matplotlib hexbin documentation](#) for more.

Pie plot

You can create a pie plot with `DataFrame.plot.pie()` or `Series.plot.pie()`. If your data includes any `NaN`, they will be automatically filled with 0. A `ValueError` will be raised if there are any negative values in your data.

[Skip to main content](#)


```
In [85]: series = pd.Series(3 * np.random.rand(4), index=["a", "b", "c", "d"], name="series")  
In [86]: series.plot.pie(figsize=(6, 6));
```

[Skip to main content](#)


For pie plots it's best to use square figures, i.e. a figure aspect ratio 1. You can create the figure with equal width and height, or force the aspect ratio to be equal after plotting by calling

`ax.set_aspect('equal')` on the returned `axes` object.

Note that pie plot with `DataFrame` requires that you either specify a target column by the `y` argument or `subplots=True`. When `y` is specified, pie plot of selected column will be drawn. If `subplots=True` is specified, pie plots for each column are drawn as subplots. A legend will be drawn in each pie plots by default; specify `legend=False` to hide it.

```
In [87]: df = pd.DataFrame(
....:     3 * np.random.rand(4, 2), index=["a", "b", "c", "d"], columns=["x", "y"]
....: )
....:

In [88]: df.plot.pie(subplots=True, figsize=(8, 4));
```



[Skip to main content](#)

You can use the `labels` and `colors` keywords to specify the labels and colors of each wedge.


Warning

Most pandas plots use the `label` and `color` arguments (note the lack of “s” on those). To be consistent with `matplotlib.pyplot.pie()` you must use `labels` and `colors`.

If you want to hide wedge labels, specify `labels=None`. If `fontsize` is specified, the value will be applied to wedge labels. Also, other keywords supported by `matplotlib.pyplot.pie()` can be used.

```
In [89]: series.plot.pie(  
....:     labels=["AA", "BB", "CC", "DD"],  
....:     colors=["r", "g", "b", "c"],  
....:     autopct="%.2f",  
....:     fontsize=20,  
....:     figsize=(6, 6),  
....: );  
....:
```


[Skip to main content](#)



If you pass values whose sum total is less than 1.0 they will be rescaled so that they sum to 1.

```
In [90]: series = pd.Series([0.1] * 4, index=["a", "b", "c", "d"], name="series2")
```

[Skip to main content](#)



See the [matplotlib pie documentation](#) for more.

Plotting with missing data

[Skip to main content](#)

pandas tries to be pragmatic about plotting `DataFrames` or `Series` that contain missing data. Missing values are dropped, left out, or filled depending on the plot type.

Plot Type	NaN Handling
Line	Leave gaps at NaNs
Line (stacked)	Fill 0's
Bar	Fill 0's
Scatter	Drop NaNs
Histogram	Drop NaNs (column-wise)
Box	Drop NaNs (column-wise)
Area	Fill 0's
KDE	Drop NaNs (column-wise)
Hexbin	Drop NaNs
Pie	Fill 0's

If any of these defaults are not what you want, or if you want to be explicit about how missing values are handled, consider using `fillna()` or `dropna()` before plotting.

Plotting tools

[Skip to main content](#)

These functions can be imported from `pandas.plotting` and take a `Series` or `DataFrame` as an argument.

Scatter matrix plot


You can create a scatter plot matrix using the `scatter_matrix` method in `pandas.plotting`:

```
In [92]: from pandas.plotting import scatter_matrix
```

```
In [93]: df = pd.DataFrame(np.random.randn(1000, 4), columns=["a", "b", "c", "d"])
```

```
In [94]: scatter_matrix(df, alpha=0.2, figsize=(6, 6), diagonal="kde");
```

[Skip to main content](#)




Density plot

[Skip to main content](#)


```
In [95]: ser = pd.Series(np.random.randn(1000))
```

```
In [96]: ser.plot.kde();
```



Andrews curves

Andrews curves allow one to plot multivariate data as a large number of curves that are created

[Skip to main content](#)

information. By coloring these curves differently for each class it is possible to visualize data clustering. Curves belonging to samples of the same class will usually be closer together and form larger structures.

Note: The “Iris” dataset is available [here](#).


```
In [97]: from pandas.plotting import andrews_curves
```

```
In [98]: data = pd.read_csv("data/iris.data")
```

```
In [99]: plt.figure();
```

```
In [100]: andrews_curves(data, "Name");
```

[Skip to main content](#)




Parallel coordinates

Parallel coordinates is a plotting technique for plotting multivariate data, see the [Wikipedia entry](#) for an introduction. Parallel coordinates allows one to see clusters in data and to estimate other statistics visually. Using parallel coordinates points are represented as connected line segments. Each vertical line represents one attribute. One set of connected line segments represents one data point. Points that tend to cluster will appear closer together.

[Skip to main content](#)

```
In [101]: from pandas.plotting import parallel_coordinates  
  
In [102]: data = pd.read_csv("data/iris.data")  
  
In [103]: plt.figure();  
  
In [104]: parallel_coordinates(data, "Name");
```




[Skip to main content](#)

Lag plots are used to check if a data set or time series is random. Random data should not exhibit any structure in the lag plot. Non-random structure implies that the underlying data are not random. The `lag` argument may be passed, and when `lag=1` the plot is essentially `data[:-1]` vs. `data[1:]`.

```
In [105]: from pandas.plotting import lag_plot
In [106]: plt.figure();
In [107]: spacing = np.linspace(-99 * np.pi, 99 * np.pi, num=1000)
In [108]: data = pd.Series(0.1 * np.random.rand(1000) + 0.9 * np.sin(spacing))
In [109]: lag_plot(data);
```

[Skip to main content](#)




Autocorrelation plot

Autocorrelation plots are often used for checking randomness in time series. This is done by computing autocorrelations for data values at varying time lags. If time series is random, such autocorrelations should be near zero for any and all time-lag separations. If time series is non-random then one or more of the autocorrelations will be significantly non-zero. The horizontal lines displayed in the plot correspond to 95% and 99% confidence bands. The dashed line is 99%

[Skip to main content](#)

```
In [110]: from pandas.plotting import autocorrelation_plot  
  
In [111]: plt.figure();  
  
In [112]: spacing = np.linspace(-9 * np.pi, 9 * np.pi, num=1000)  
  
In [113]: data = pd.Series(0.7 * np.random.rand(1000) + 0.3 * np.sin(spacing))  
  
In [114]: autocorrelation_plot(data);
```



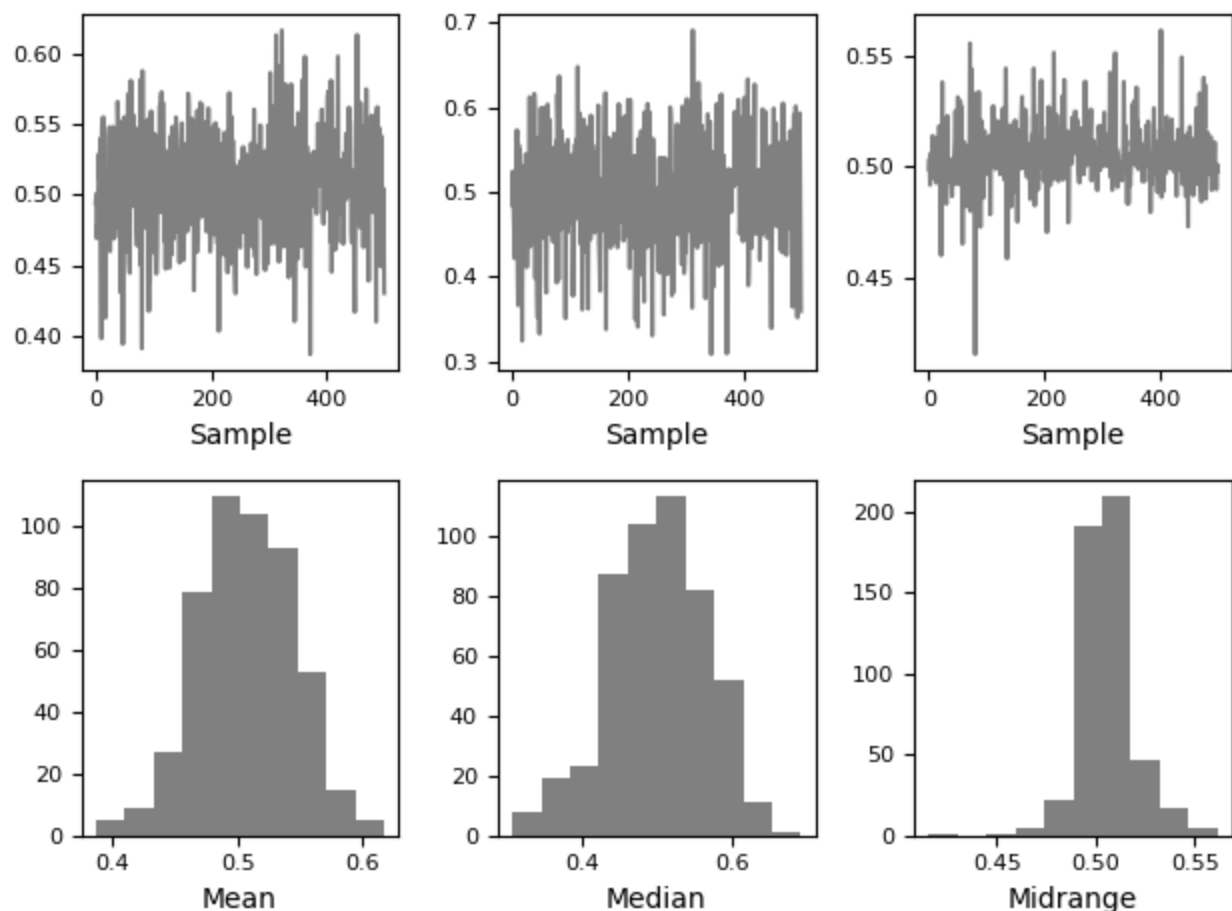
[Skip to main content](#)

Bootstrap plot

Bootstrap plots are used to visually assess the uncertainty of a statistic, such as mean, median, midrange, etc. A random subset of a specified size is selected from a data set, the statistic in question is computed for this subset and the process is repeated a specified number of times. Resulting plots and histograms are what constitutes the bootstrap plot.

```
In [115]: from pandas.plotting import bootstrap_plot  
  
In [116]: data = pd.Series(np.random.rand(1000))  
  
In [117]: bootstrap_plot(data, size=50, samples=500, color="grey");
```

[Skip to main content](#)



RadViz

RadViz is a way of visualizing multi-variate data. It is based on a simple spring tension minimization algorithm. Basically you set up a bunch of points in a plane. In our case they are equally spaced on a unit circle. Each point represents a single attribute. You then pretend that each sample in the data set is attached to each of these points by a spring, the stiffness of which is proportional to the numerical value of that attribute (they are normalized to unit interval). The point in the plane, where


[Skip to main content](#)

representing our sample will be drawn. Depending on which class that sample belongs it will be colored differently. See the R package [Radviz](#) for more information.

Note: The "Iris" dataset is available [here](#).

```
In [118]: from pandas.plotting import radviz  
  
In [119]: data = pd.read_csv("data/iris.data")  
  
In [120]: plt.figure();  
  
In [121]: radviz(data, "Name");
```

[Skip to main content](#)



Plot formatting

Setting the plot style

From version 1.5 and up, matplotlib offers a range of pre-configured plotting styles. Setting the style can be used to easily give plots the general look that you want. Setting the style is as easy as

[Skip to main content](#)

calling `matplotlib.style.use(my_plot_style)` before creating your plot. For example you could write `matplotlib.style.use('ggplot')` for ggplot-style plots.

You can see the various available style names at `matplotlib.style.available` and it's very easy to try them out.

General plot style arguments

Most plotting methods have a set of keyword arguments that control the layout and formatting of the returned plot:

```
In [122]: plt.figure();
```

```
In [123]: ts.plot(style="k--", label="Series");
```

[Skip to main content](#)



For each kind of plot (e.g. `line`, `bar`, `scatter`) any additional arguments keywords are passed along to the corresponding matplotlib function (`ax.plot()`, `ax.bar()`, `ax.scatter()`). These can be used to control additional styling, beyond what pandas provides.

Controlling the legend

You may set the `legend` argument to `False` to hide the legend, which is shown by default.

[Skip to main content](#)

```
In [124]: df = pd.DataFrame(np.random.randn(1000, 4), index=ts.index, columns=list("ABCD"))  
  
In [125]: df = df.cumsum()  
  
In [126]: df.plot(legend=False);
```




Controlling the labels

[Skip to main content](#)

You may set the `xlabel` and `ylabel` arguments to give the plot custom labels for x and y axis. By default, pandas will pick up index name as xlabel, while leaving it empty for ylabel.

```
In [127]: df.plot();
```

```
In [128]: df.plot(xlabel="new x", ylabel="new y");
```




[Skip to main content](#)

Scales

You may pass `logy` to get a log-scale Y axis.

```
In [129]: ts = pd.Series(np.random.randn(1000), index=pd.date_range("1/1/2000", periods=1000))  
In [130]: ts = np.exp(ts.cumsum())  
In [131]: ts.plot(logy=True);
```

[Skip to main content](#)




See also the `logx` and `loglog` keyword arguments.

Plotting on a secondary y-axis

To plot data on a secondary y-axis, use the `secondary_y` keyword:

```
In [132]: df["A"].plot();
```


[Skip to main content](#)



To plot some columns in a `DataFrame`, give the column names to the `secondary_y` keyword:

```
In [134]: plt.figure();  
  
In [135]: ax = df.plot(secondary_y=["A", "B"])  
  
In [136]: ax.set_ylabel("CD scale");  
  
In [137]: ax.right_ax.set_ylabel("AB scale");
```

[Skip to main content](#)




Note that the columns plotted on the secondary y-axis is automatically marked with "(right)" in the legend. To turn off the automatic marking, use the `mark_right=False` keyword:

```
In [138]: plt.figure();
```

```
In [139]: df.plot(secondary_y=["A", "B"], mark_right=False);
```

[Skip to main content](#)



Custom formatters for timeseries plots

pandas provides custom formatters for timeseries plots. These change the formatting of the axis labels for dates and times. By default, the custom formatters are applied only to plots created by pandas with `DataFrame.plot()` or `Series.plot()`. To have them apply to all plots, including those made by matplotlib, set the option `pd.options.plotting.matplotlib.register_converters = True` or use `pandas.plotting.register_matplotlib_converters()`

[Skip to main content](#)


Suppressing tick resolution adjustment

pandas includes automatic tick resolution adjustment for regular frequency time-series data. For limited cases where pandas cannot infer the frequency information (e.g., in an externally created `twinx`), you can choose to suppress this behavior for alignment purposes.

Here is the default behavior, notice how the x-axis tick labeling is performed:

```
In [140]: plt.figure();  
In [141]: df["A"].plot();
```

[Skip to main content](#)




Using the `x_compat` parameter, you can suppress this behavior:

```
In [142]: plt.figure();
```

```
In [143]: df["A"].plot(x_compat=True);
```

[Skip to main content](#)




If you have more than one plot that needs to be suppressed, the `use` method in `pandas.plotting.plot_params` can be used in a `with` statement:

```
In [144]: plt.figure();
```

```
In [145]: with pd.plotting.plot_params.use("x_compat", True):
.....:     df["A"].plot(color="r")
.....:     df["B"].plot(color="g")
.....:     df["C"].plot(color="b")
.....:
```

[Skip to main content](#)



Automatic date tick adjustment

`TimedeltaIndex` now uses the native matplotlib tick locator methods, it is useful to call the automatic date tick adjustment from matplotlib for figures whose ticklabels overlap.

See the `autofmt_xdate` method and the [matplotlib documentation](#) for more.


[Skip to main content](#)

Subplots

Each `Series` in a `DataFrame` can be plotted on a different axis with the `subplots` keyword:

```
In [146]: df.plot(subplots=True, figsize=(6, 6));
```

[Skip to main content](#)



Using layout and targeting multiple axes

[Skip to main content](#)

The layout of subplots can be specified by the `layout` keyword. It can accept `(rows, columns)`. The `layout` keyword can be used in `hist` and `boxplot` also. If the input is invalid, a `ValueError` will be raised.


The number of axes which can be contained by rows x columns specified by `layout` must be larger than the number of required subplots. If layout can contain more axes than required, blank axes are not drawn. Similar to a NumPy array's `reshape` method, you can use `-1` for one dimension to automatically calculate the number of rows or columns needed, given the other.

```
In [147]: df.plot(subplots=True, layout=(2, 3), figsize=(6, 6), sharex=False);
```

[Skip to main content](#)

[Getting started](#)[User Guide](#)[API reference](#)[Development](#)[Release notes](#)

2.1.1

[Intro to data structures](#)[Essential basic functionality](#)[IO tools \(text, CSV, HDF5, ...\)](#)[PyArrow Functionality](#)[Indexing and selecting data](#)[MultiIndex / advanced indexing](#)[Copy-on-Write \(CoW\)](#)[Merge, join, concatenate and compare](#)[Reshaping and pivot tables](#)[Working with text data](#)[Working with missing data](#)[Duplicate Labels](#)[Categorical data](#)[Nullable integer data type](#)[Nullable Boolean data type](#)[Chart visualization](#)[Table Visualization](#)

The above example is identical to using:

```
In [148]: df.plot(subplots=True, layout=(2, -1), figsize=(6, 6), sharex=False);
```

[Skip to main content](#)

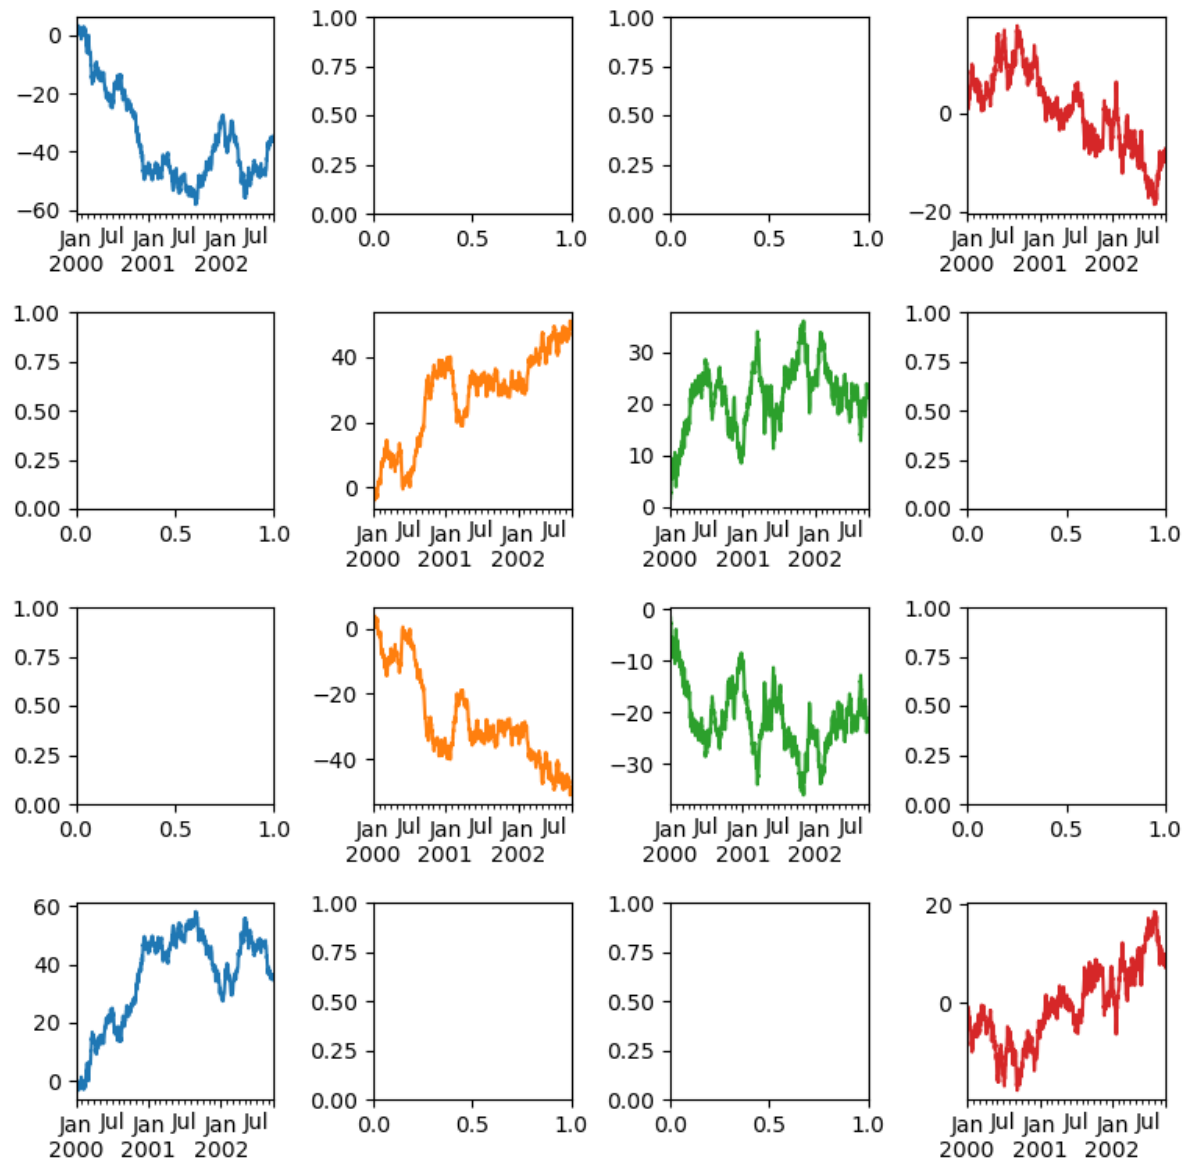
The required number of columns (3) is inferred from the number of series to plot and the given number of rows (2).

You can pass multiple axes created beforehand as list-like via `ax` keyword. This allows more complicated layouts. The passed axes must be the same number as the subplots being drawn.

When multiple axes are passed via the `ax` keyword, `layout`, `sharex` and `sharey` keywords don't affect to the output. You should explicitly pass `sharex=False` and `sharey=False`, otherwise you will see a warning.

```
In [149]: fig, axes = plt.subplots(4, 4, figsize=(9, 9))
In [150]: plt.subplots_adjust(wspace=0.5, hspace=0.5)
In [151]: target1 = [axes[0][0], axes[1][1], axes[2][2], axes[3][3]]
In [152]: target2 = [axes[3][0], axes[2][1], axes[1][2], axes[0][3]]
In [153]: df.plot(subplots=True, ax=target1, legend=False, sharex=False, sharey=False);
In [154]: (-df).plot(subplots=True, ax=target2, legend=False, sharex=False, sharey=False);
```

[Skip to main content](#)

[Skip to main content](#)

```
In [155]: np.random.seed(123456)

In [156]: ts = pd.Series(np.random.randn(1000), index=pd.date_range("1/1/2000", periods=1000))

In [157]: ts = ts.cumsum()

In [158]: df = pd.DataFrame(np.random.randn(1000, 4), index=ts.index, columns=list("ABCD"))

In [159]: df = df.cumsum()
```

```
In [160]: fig, axes = plt.subplots(nrows=2, ncols=2)

In [161]: plt.subplots_adjust(wspace=0.2, hspace=0.5)

In [162]: df["A"].plot(ax=axes[0, 0]);

In [163]: axes[0, 0].set_title("A");

In [164]: df["B"].plot(ax=axes[0, 1]);

In [165]: axes[0, 1].set_title("B");


In [166]: df["C"].plot(ax=axes[1, 0]);

In [167]: axes[1, 0].set_title("C");

In [168]: df["D"].plot(ax=axes[1, 1]);

In [169]: axes[1, 1].set_title("D");
```

[Skip to main content](#)



Plotting with error bars

Plotting with error bars is supported in `DataFrame.plot()` and `Series.plot()`.

Horizontal and vertical error bars can be supplied to the `xerr` and `yerr` keyword arguments to `plot()`. The error values can be specified using a variety of formats:

[Skip to main content](#)

- As a `DataFrame` or `dict` of errors with column names matching the `columns` attribute of the plotting `DataFrame` or matching the `name` attribute of the `Series`.
- As a `str` indicating which of the columns of plotting `DataFrame` contain the error values.
- As raw values (`list`, `tuple`, or `np.ndarray`). Must be the same length as the plotting `DataFrame` / `Series`.

Here is an example of one way to easily plot group means with standard deviations from the raw data.

```
# Generate the data
In [170]: ix3 = pd.MultiIndex.from_arrays(
.....:     [
.....:         ["a", "a", "a", "a", "a", "b", "b", "b", "b", "b"],
.....:         ["foo", "foo", "foo", "bar", "bar", "foo", "foo", "bar", "bar", "bar"],
.....:     ],
.....:     names=["letter", "word"],
.....: )

In [171]: df3 = pd.DataFrame(
.....:     {
.....:         "data1": [9, 3, 2, 4, 3, 2, 4, 6, 3, 2],
.....:         "data2": [9, 6, 5, 7, 5, 4, 5, 6, 5, 1],
.....:     },
.....:     index=ix3,
.....: )

# Group by index labels and take the means and standard deviations
# for each group
In [172]: gp3 = df3.groupby(level=("letter", "word"))

In [173]: means = gp3.mean()

In [174]: errors = gp3.std()
```

[Skip to main content](#)

Out[175]:

		data1	data2
letter	word		
a	bar	3.500000	6.000000
	foo	4.666667	6.666667
b	bar	3.666667	4.000000
	foo	3.000000	4.500000

In [176]: errors

Out[176]:


		data1	data2
letter	word		
a	bar	0.707107	1.414214
	foo	3.785939	2.081666
b	bar	2.081666	2.645751
	foo	1.414214	0.707107

Plot

In [177]: fig, ax = plt.subplots()

In [178]: means.plot.bar(yerr=errors, ax=ax, capsize=4, rot=0);

[Skip to main content](#)



Asymmetrical error bars are also supported, however raw error values must be provided in this

© 2023 pandas via NumFOCUS, Inc. Hosted by OVHcloud.

Built with the PyData Sphinx Theme 0.13.3.

Created using Sphinx 6.2.1.